

University of Bremen

TODO: Make sure
this is centered!

figures/unibremen.pdf

Faculty 3 — Mathematics & Computer Science

Master's Thesis

Building Code Cities using the Language Server Protocol — *DRAFT* —

Falko Galperin

- | | |
|--------------------|---|
| 1. <i>Reviewer</i> | Prof. Dr. Rainer Koschke
<i>Working group Software Engineering</i> |
| 2. <i>Reviewer</i> | Prof. Dr. Ute Bormann
<i>Working group Computer Networks</i> |

October 6, 2024

Abstract

TODO!

Declaration

I hereby declare that I have completed this master's thesis independently and without any external assistance, unless explicitly stated otherwise. I have not used any sources or aids other than those specified. All passages that have been quoted verbatim or paraphrased from published sources are clearly identified as such.

Bremen, October 6, 2024

Falko Galperin

Acknowledgements

TODO!


Contents

1	Introduction	1
1.1	Format	1
1.2	Motivation	2
1.2.1	Code Cities & SEE	2
1.2.2	Language Server Protocol	4
1.2.3	Integration	7
1.3	Goals & Research Questions	7
1.4	Thesis Structure	9
2	Concepts	11
2.1	SEE	11
2.1.1	Basics	12
2.1.2	Project Graph	13
2.1.3	Other Relevant Features	14
2.2	Language Server Protocol	15
2.2.1	Basics	16
2.2.2	Capabilities	17
2.3	Interim Conclusion	22
3	Implementation	25
3.1	Preliminary Changes	25
3.1.1	Specification Cleanup	25
3.1.2	Preparing SEE	25
3.2	Generating Code Cities using LSP	25
3.3	Integrating LSP Functionality into Code Cities	25
3.4	Integrating LSP Functionality into Code Windows	26
3.5	Technical Evaluation	26
3.6	Interim Conclusion	26
4	Evaluation	27
4.1	Plan	27
4.2	Structure	27

4.2.1	Questionnaire	27
4.2.2	Tasks	27
4.3	Results	27
4.4	Threats to Validity	28
4.5	Interim Conclusion	28
5	Conclusion	29
5.1	Limitations	29
5.2	Future Work	29
5.3	The End	29
A	List of TODOs	31
B	Glossary	35
C	Acronyms	37
D	Attached Files	39
E	List of Figures	41
F	List of Tables	43
	List of Listings	45
G	List of Listings	45
H	Bibliography	47

1

Introduction

 HIS master's thesis is about the integration of the Language Server Protocol into code cities (specifically, into SEE). The main novel contribution will be a way of creating code cities using only information provided by the Language Server Protocol, while additional contributions consist of the integration of Language Server Protocol-based functionality into code cities as well as their integration into SEE's code windows. Finally, the penultimate chapter 4 describes a controlled experiment (with $n = 4$ participants) in which code cities are compared to traditional IDEs via a user study—this serves as an evaluation for this thesis and for code cities as IDE replacements in general.

In this chapter, apart from explaining some formatting semantics, we will examine the motivation and basics behind each of the central concepts (i. e., code cities, the Language Server Protocol, and the integration of the two), name the goals and research question of the thesis, and finally describe the structure of upcoming chapters.

1.1 Format

This document uses many technical terms that not every reader may know. To remedy this, starting from the next section, whenever a technical term or acronym appears for the first time, it will be printed in *this color*, and an explanation of that term will appear in a box of the same color nearby. Terms that are already explained in the text itself will not receive such a box. The terms and acronyms are collected within the glossaries in appendices B and C—all mentions of such terms also link (in the digital version of this thesis, at least) to the corresponding part of the glossary. In total, the following colors are used to convey specific meanings:

- **Maroon** for the introduction of a glossary term or acronym,
- **Fuchsia** for internal links (e. g., to other sections),

New Term: *LOC* (Lines of Code)

The number of lines in a source code file.

- **Blue** for external links (e. g., to web pages),
- **Green** for cited literature, and
- **Cyan** for references to attached files (see appendix D).

1.2 Motivation

As mentioned at the beginning of this section, this master’s thesis is about *integrating* the *Language Server Protocol* into *Code Cities*. I will motivate each of these italicized central points individually in the following sections. Note that these will get more thorough explanations in chapter 2.

1.2.1 Code Cities & SEE

Visualization in general often helps facilitate the understanding of complex systems by representing them with a simplified visual model. This can be especially useful in the area of software engineering, where it is often hard to get an intuitive overview of large software systems. One such software visualization—called *Software Engineering Experience* (**SEE**)—is being developed at the University of Bremen and will be introduced in the next section.

SEE is an interactive software visualization tool using the *code city* metaphor in 3D, developed in the Unity game engine. It features collaborative “multiplayer” functionality across multiple platforms¹, allowing multiple participants to view and interact with the same code city together.

In the code city metaphor, software components are visualized as buildings within a city. Various metrics from the original software can then be represented by different visual properties of each building—for example, the *Lines of Code* (**LOC**) within a file might correlate to the height of the corresponding building. Relationships between software components, such as where components are referenced, are instead represented by edges drawn between the respective buildings. The exception to this are part-of relations, that is, relations that describe which component belongs to which other component. These are

¹Notably, besides usual desktop and touchscreen-based environments, virtual reality (e.g., via the *Valve Index*) is supported as well.

figures/SEE_screenshot.png

Figure 1.1: A code city visualized in SEE.

New Term: GXL (Graph eXchange Language)

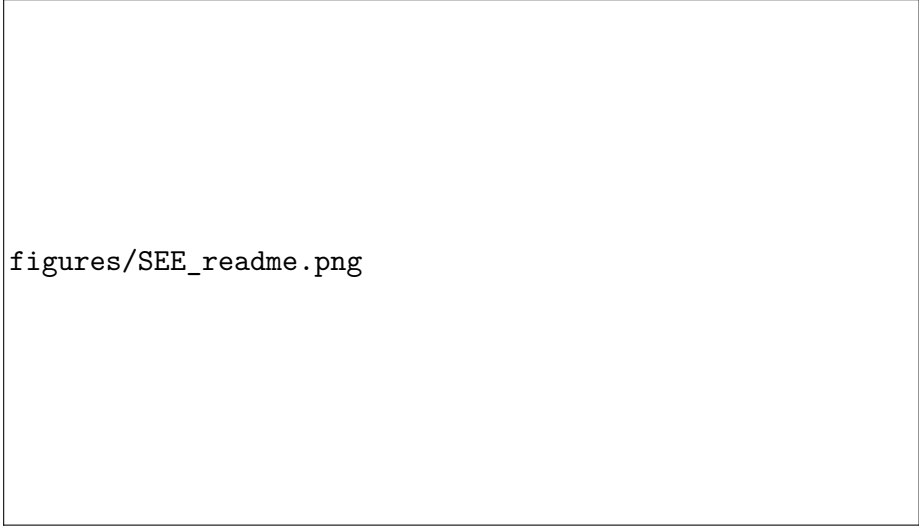
A file format for graphs, used in SEE for representing dependency and hierarchy graphs of software projects.

instead visualized in SEE by buildings being nested within their corresponding “parent” building. In this way, the data model of SEE can be represented as a graph in which the software components are the nodes and the relationships are the edges.

For example, in fig. 1.1, we can see the source code of the SPOTBUGS project² rendered as a code city. A few very tall buildings—indicating that the respective component is very big and that a refactoring into smaller pieces may be in order—immediately jump out. Additionally, this visualization also makes the number of methods readily apparent: the redder a node, the higher its method count. fig. 1.2 instead visualizes the modeled architecture of a very small system, as compared to a city “empirically” generated by an implementation like in the previous example. Here, we can also see yellow edges between the components, in this case representing desired references that should be present between components.

Currently, code cities in SEE are rendered by reading in pre-made *Graph eXchange Language* (GXL) files, which can be created by the proprietary Axivion Suite. This approach has the disadvantage of only supporting languages supported by the Axivion Suite, as well as making regenerating cities (e.g., if the source code changed) fairly cumbersome. Another current shortcoming of SEE is that information about the source code available to the user is limited when compared to an *Integrated Development Environment* (IDE)—for

²<https://github.com/spotbugs/spotbugs> (last access: 2024-09-11)



figures/SEE_readme.png

Figure 1.2: An example of what edges can look like in SEE. **TODO:** Use higher quality image.

New Term: *IDE* (Integrated Development Environment)

Editor for source code with features that are useful for development (e. g., highlighting errors). Examples are *Eclipse* or *JETBRAINS IntelliJ*.

example, quickly displaying documentation for a given component by hovering over it is not supported. This is where the Language Server Protocol can help.

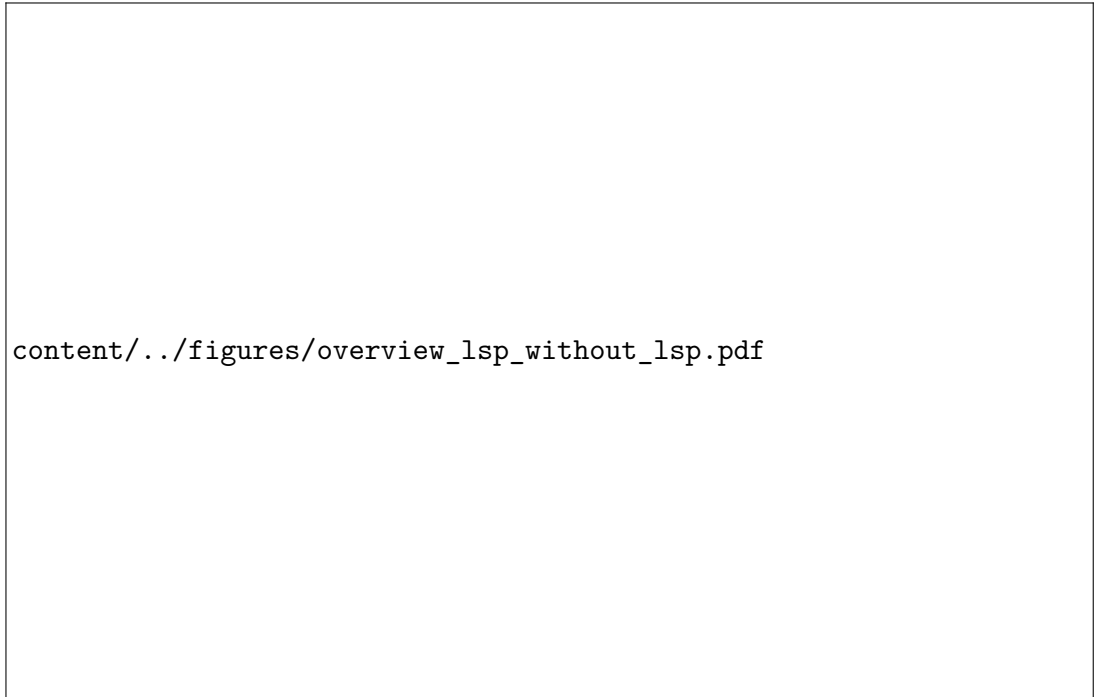
1.2.2 Language Server Protocol

As stated on its website:

Adding features like auto complete, go to definition, or documentation on hover for a programming language takes significant effort. Traditionally[,] this work had to be repeated for each development tool, as each tool provides different APIs for implementing the same feature.

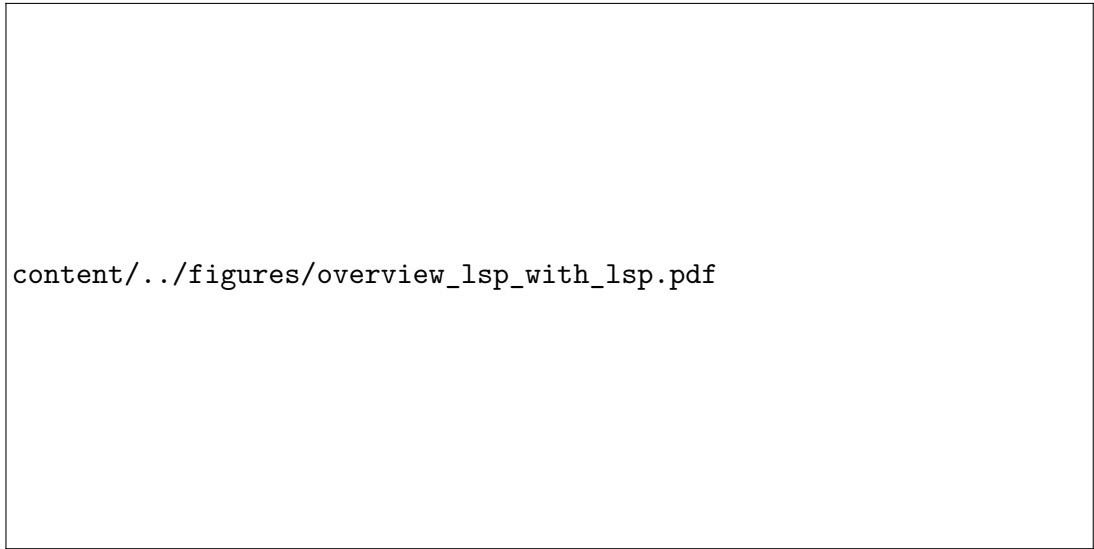
A *Language Server* is meant to provide the language-specific smarts and communicate with development tools over a protocol that enables inter-process communication.

The idea behind the *Language Server Protocol* (LSP) is to standardize the protocol for how such servers and development tools communicate. This way, a single Language Server can be re-used in multiple development tools, which in turn can support multiple languages with minimal effort. ([LspSpec])



content/../../figures/overview_lsp_without_lsp.pdf

(a) IDE development without LSP.



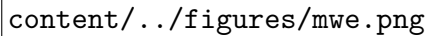
content/../../figures/overview_lsp_with_lsp.pdf

(b) IDE development with LSP.

Figure 1.3: An illustration of how LSP can help simplify IDE development.

New Term: *Language Client*

A development tool, such as an IDE, that supports LSP and can hence integrate language-specific features into itself using compatible Language Servers.



```
content/../../figures/mwe.png
```

Figure 1.4: An example of the texlab Language Server running in NEOVIM.

Since LSP³ is a central component of my master’s thesis, I have created a diagram in fig. 1.3 in the hope to strengthen intuitions around the motivation and use of the protocol. While the LSP specification has originally been created by Microsoft, it is by now an open-source project⁴, where changes can be actively proposed using issues or pull requests. Apart from the specification itself, a great number of open-source implementations of Language Servers for all kinds of programming languages from Ada to Zig exist. A partial overview of available implementations is listed at <https://microsoft.github.io/language-server-protocol/implementors/servers/> (*last access: 2024-09-11*).

The protocol introduces the concept of so-called *capabilities*, which define a specific set of features a given Language Server (and *Language Client*) support. These include navigational features, like the ability to jump to a variable’s declaration, and editing-related features, such as autocomplete. To give a specific example of what an LSP capability might look like in practice, the texlab⁵ Language Server for \LaTeX —which I am using while writing this document—provides a list of available packages when one starts typing text after “\usepackage{”. Additionally, for the currently hovered package, a short description of it is displayed. A screenshot of this behavior within the NEOVIM editor is provided in fig. 1.4.

The counterparts to Language Servers are the Language Clients: These are the IDEs and editors that incorporate the Language Server into themselves. Examples for IDEs that support acting as a Language Client in the LSP context include *Eclipse*, *Emacs*, JETBRAINS IntelliJ (NEO)VIM, and *Visual Studio Code*.

³Note that I will often refer to the Language Server Protocol as just “LSP” instead of “the LSP” (e.g., “IDEs use LSP”) from now on, as this is how the specification [LspSpec] does it as well.

⁴Available at <https://github.com/Microsoft/language-server-protocol> (*last access: 2024-09-11*).

⁵<https://github.com/latex-lsp/texlab> (*last access: 2024-09-12*)

1.2.3 Integration

TODO: Motivate integration

1.3 Goals & Research Questions

The goal of this master’s thesis—as outlined in section 1.2.3—is to integrate the Language Server Protocol into SEE by making it a Language Client, then evaluate this implementation by comparing it with traditional IDEs in a user study. To this end, the main contribution is a way of generating code cities using the Language Server, where all the information obtainable by relevant LSP capabilities should be manifested⁶ in the city in a suitable way. This is an unintended (or at the very least, unusual) use of LSP and may require some experimentation.

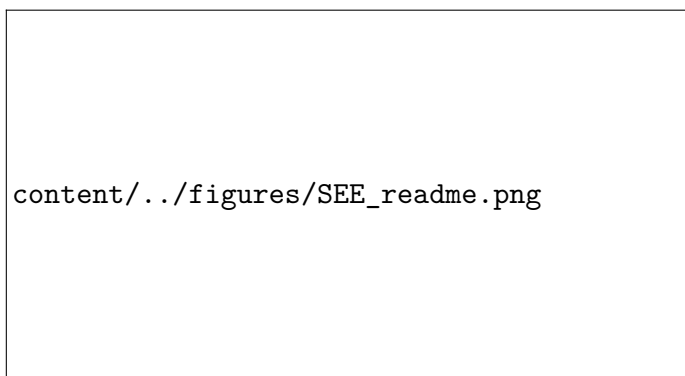


Figure 1.6: A code window.

Apart from the code cities, SEE also provides the so-called *code windows*, in which the source code of a specific component can be viewed in a similar way as in an IDE. This can be seen in fig. 1.6. An additional goal of this master’s thesis is to enhance the functionality of the code windows by implementing more IDE-like behavior into it (e. g., allowing users to go to

a variable’s declaration, or displaying diagnostics inline) using the Language Server.


It should be noted that any LSP capabilities which involve modifying the underlying software project is out of scope for this master’s thesis. Rewriting the code windows to be editable (in a way that distributes the edits over the network) is complex enough to warrant its own thesis [see also Ble22], not even taking into account that this would also more than double the number of capabilities that would then become useful to implement. As such, only capabilities that are “read-only” (i. e., that do not modify the source code or project structure in any way) will be considered when planning which features to implement as part of the thesis. Consult ?? for a full list of LSP capabilities which I plan to implement.

Additionally, I will not implement the C# interface to the LSP (i. e., translating between C# method calls and *JavaScript Object Notation—Remote Procedure Call* (JSON-RPC) calls).

⁶Since there is a lot of diverse data available via LSP, it makes sense to only immediately display the most pertinent information and make the rest of it available upon request within SEE’s UI.

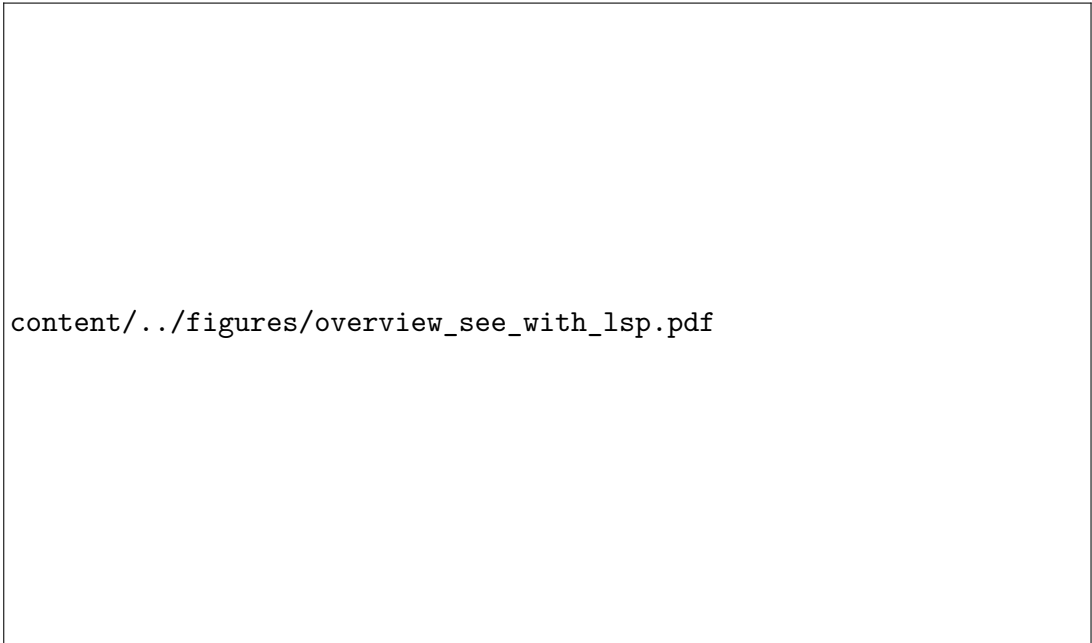
TODO: Maybe remove the next part if already mentioned before.

TODO: In general, reword some stuff from “planning to do” ⇒ “done”



`content/../../figures/overview_see_without_lsp.pdf`

(a) SEE as it exists right now. The analyzers on the right are developed by Axivion.



`content/../../figures/overview_see_with_lsp.pdf`

(b) SEE with implemented LSP integration. The Language Servers on the right are mostly developed by the open-source community.

Figure 1.5: A simplified illustration of how SEE currently gains and uses information about software projects, and how the integration of LSP would change that.

New Term: *JSON-RPC* (JavaScript Object Notation—Remote Procedure Call)

A remote procedure call protocol that uses JSON as its encoding, supporting (among other features) asynchronous calls and notifications. It is used as the base for LSP (even though LSP is technically not a remote protocol).

There already exist well-made interfaces for this purpose⁷, and the focus of the thesis will be on the integration of the protocol's *data* into SEE's code cities and code windows, not the integration of the protocol itself.

Hence, the goals for this thesis can be summarized as follows:

- Integrating an LSP framework into SEE and allowing users to manage Language Servers from within SEE
- Making SEE a Language Client, such that:
 1. Code cities can be generated directly from source code directories, using Language Servers that SEE interfaces with,
 2. code windows gain “read-only” IDE-like functionality, covering behavior of capabilities listed in ??, and
 3. Code cities gain similar functionality (where applicable), such as displaying relevant documentation when hovering over a node.
- Evaluating the above empirically in a controlled experiment via a user study.

The main research questions that I want to answer in this thesis are as follows:

RQ1 Is it feasible to generate code cities using the Language Server Protocol?

RQ2 Are code cities a suitable means to present LSP information to developers as compared to IDEs + tables (on the dimensions of speed, accuracy, and usability)?

TODO: This is rather vague. Is that alright or do I need to operationalize here?

1.4 Thesis Structure

We begin by examining the Language Server Protocol and the concept of code cities along with SEE more closely in chapter 2. Next, in chapter 3 we take a look at my implementation of the LSP-based code city generation algorithm, alongside additional contributions to visualize this information in the cities and code windows of SEE. In the course of this, we will answer RQ1. To answer RQ2, I will carry out a user study comparing an LSP-enabled

⁷Such as OmniSharp's implementation of LSP in C#, which I plan to use: <https://github.com/OmniSharp/csharp-language-server-protocol> (last access: 2024-28-01)

IDE (specifically, VSCode) with an LSP-enabled code city visualization (specifically, SEE) and report on its results in chapter 4. Finally, we wrap up in chapter 5 by summarizing the thesis's results and providing an outlook on additional ideas for the implementation, while also listing some possible avenues of further research.

Concepts



BEFORE tackling the implementation, we need to take a closer look at the concepts central to this thesis, so that we can form a concrete idea of which parts of LSP are well-suited to being integrated into code cities. Thus, we will examine the concept of code cities, where we will use SEE as a concrete implementation (which we do in section 2.1), as well as the Language Server Protocol (which we do in section 2.2). For the former, we will go over some of the existing literature regarding code cities (as this is more of an academic topic than LSP) and describe the essentials that we need to work with in the implementation, such as the project graph. As for the latter, we will go over the available capabilities and take a look at both existing Language Servers and Language Clients to get an idea of what the Language Server Protocol offers and how it is most commonly used. For both topics, we will focus on the parts relevant to the implementation and evaluation—for example, we will only explore those capabilities in detail that actually end up being used in this thesis.

2.1 SEE

As explained in section 1.2.1, SEE is an interactive software visualization tool using the code city metaphor in 3D, with the aim to make it easier to work with large software projects on an architectural level. To name a few example scenarios in which using code cities could be useful, **compared to** using traditional IDEs:

- Senior developers may use it in its “multiplayer” mode to explain the structure of their software to newcomers
- Project planners could visualize *code smells* to find candidates for refactoring [Gal21]
- Software architects can find violations of the planned architecture using SEE’s *reflexion analysis*

TODO: Back this up with sources

TODO: Refer back here later, noting how LSP can generate code cities for each of these use cases

New Term: *Code Smell*

Certain structures in source code that suggest that a refactoring is in order, such as duplicated code, or a very long method [FB19, pp. 85–87].

New Term: *Reflexion Analysis*

The process of comparing the architecture and implementation of a software project and finding incongruencies between the two.

SEE is an open-source project, currently hosted at GitHub¹. It is a research project at the University of Bremen, where it has been in development **since 2019**, and where it is frequently offered as a bachelor or master project for the students. While it was at first a project for the *Unreal Engine*, the game engine has been switched to *Unity* relatively early, mainly because its editor eased development due to the ability to reload the *User Interface* (UI) without having to restart the whole engine.

TODO: is this correct? And put source here

After going over the basics of how SEE works, I will give an explanation and formalization of the project graph—as this is the central component in the LSP-based code city-generation algorithm—followed by a brief overview over other features that are relevant for this thesis.

2.1.1 Basics

As mentioned before, when analyzing a software project in SEE, its source code structure is reflected in the rendered buildings, where connections between the components, such as references, are displayed using edges rendered as splines (this will be elaborated in section 2.1.2). Metrics of the project, such as the McCabe complexity [McC76], can then be mapped onto visual properties of this rendering, such as its color, width, or height.

TODO: Needs a 3D SEE screenshot, esp.

So far, the kind of visualization described here would still be possible **in two dimensions**, but adding a third dimension gives us the benefit of having another axis to map metrics onto, as well as making it a more natural environment for humans to interact with. By giving users the ability to walk, move the camera around, and so on, as is possible in many video games, the code city can be navigated in a richer and more intuitive manner than if it were just a simple 2D image. This also makes it possible to offer SEE as a virtual reality application. Virtual reality gives users an even more immersive and natural way of interacting with the environment, leading to improved spatial memory of the city, as some recent studies have shown. It also makes it a more fitting environment for multiple people. Each player

for edges

TODO: add 2D screenshot

TODO: Source for 3D/game/VR being easier to navigate/understand

¹<https://github.com/uni-bremen-agst/SEE> (last access: 2024-09-18) (but note that, to clone this repository, you additionally need access to the Git LFS counterpart hosted at the University of Bremen's GitLab, as this is where paid plugins for SEE are hosted.)

New Term: *Polytree*

A directed acyclic graph which also has no undirected cycles.

can have their own avatar placed in the scene, interacting with one another and the city, communicating via voice chat.

I heavily recommend readers take a look at [this introductory video](#), as an interactive game-like project like this comes across much more clearly in a video than in mere texts and disconnected snapshots.

TODO: link the video

2.1.2 Project Graph

We will now take a look at the graph representing the source structure of the analyzed software projects in SEE. One of the main goals of this thesis is generating such a graph using LSP, as opposed to the currently available methods of generating this graph, which require access to the proprietary Axivion Suite to create the GXL files used in SEE.

These graphs can be formalized as $G = (V, E, a, s, t, \ell)$, with:

- V being a set of nodes and E being a set of edges.
- $a : (V \times \mathcal{A}_K) \rightarrow \mathcal{A}_V$ associating nodes and an attribute name ($\in \mathcal{A}_K$) with a value ($\in \mathcal{A}_V$). Note that this is a partial function.
- $s : E \rightarrow V$ denoting the source node of an edge.
- $t : E \rightarrow V$ denoting the target node of an edge.
- $\ell : E \rightarrow \Sigma$ providing a label for an edge over some alphabet Σ . Also, $\text{partOf} \in \Sigma$ such that the subgraph $(V, \{e \in E \mid \ell(e) = \text{partOf}\}, a, s, t, \ell)$ is a *polytree*.

Explained in natural language: The graph's nodes (V) can be connected by directed edges (E)—these are not necessarily unique for each possible tuple of nodes, meaning that there can be more than one edge between the same two nodes. A node represents a component in the source code (e. g., a class), while edges represent relationships between them (e. g., inheritance). This is an attributed graph. Specifically, each node can have multiple attributes (distinguished by attribute keys), while each edge always has one label that denotes the type of relationship it represents. The label `partOf` additionally has a special meaning: edges with this type induce the hierarchy of the source code (e. g., classes contained in files). Hence, if we look at the graph that removes all edges except for those with label `partOf`, and we make these edges undirected, we get a tree.

TODO: A lot of “represents” here

To illustrate, a few examples of attributes are:

- **Source.Path:** the path to the file the element is contained in.
- **Source.Name:** the name of the element within the source code.
- **Type:** the type of the element (e. g., “method”).

Another attribute that is also relevant for LSP is the **Source.Range**, which is often mentioned in the upcoming section 2.2 and chapter 3. It describes a contiguous portion of source code. We formally define the domain of ranges \mathcal{R} as the Cartesian square of positions \mathcal{P}^2 , whereas the domain of positions \mathcal{P} is defined as the Cartesian square of natural numbers (including zero), that is, $\mathcal{P} = \mathbb{N}_0^2$. Hence, as a whole, $\mathcal{R} = \mathcal{P}^2 = \mathbb{N}_0^4$. Semantically, a position $(l, c) \in \mathcal{P}$ describes a zero-indexed line l and a zero-indexed character offset c , relative to the beginning of the line. A range $(s, e) \in \mathcal{R}$ can be understood as a starting position s (inclusive) and an ending position e (exclusive). We will also occasionally “unwrap” those positions and refer to a range r as $r = (s_l, s_c, e_l, e_c)$. For example, the interval of lines that the range (partially or completely) covers is then given by $[s_l, e_l)$.

I should note that the model presented here is a simplification of SEE’s actual data model for source code graphs. For example, in reality, edges can have multiple attributes, there can be multiple edge types other than `partOf` that have the polytree property, and so on. This is just what is needed to understand the LSP-based city generation algorithm presented in section 3.2.

TODO: Give examples of software projects here to illustrate graph usage?

2.1.3 Other Relevant Features

Apart from the project graph, there are a few other features of SEE we need to go over, as they become relevant when integrating LSP into code cities (in section 3.3) and code windows (in section 3.4).

Context Menu When right-clicking any node or edge, a menu opens with several context-dependent options. These include the option to delete the element, to highlight it within the code city, to open its corresponding code window, and others. We will expand this menu with LSP-specific actions in the course of section 3.3.

TODO: Show screenshot of context menu (pre-LSP)

City Editor To actually generate a code city—assuming one has a GXL file for this purpose—a customized UI component within the *Unity Editor* exists in SEE. Here, a variety of options can be configured, such as the layout of the city, the mapping of metrics to visual attributes, or the *graph providers*, which create a project graph based on optional input parameters (such as the aforementioned GXL file). After implementing the city-generation algorithm for LSP, a new graph provider with its own Unity Editor-UI shall be implemented as well.

TODO: Provide screenshot of editor

New Term: *Unity Editor*

The main UI of the Unity game engine, in which scenes can be set up, components can be configured, the game itself can be run, etc. Note that it is only used for development purposes, and hence not included within generated builds of a game.

New Term: *VSCode (Visual Studio Code)*

A proprietary, but free IDE developed by Microsoft with a plugin system from which LSP originated. See <https://code.visualstudio.com> (*last access: 2024-10-05*)

Code Windows As explained in section 1.3, there is also the option of opening code windows to view the source code of a component, an example of which is shown in fig. 1.6. Currently, these windows do little more than lexer-based syntax highlighting, so a goal is to include more IDE-like behavior by using functionality offered by LSP.

Erosion icons Following my bachelor's thesis [Gal21], SEE offers the possibility to indicate the number of code smells per component using the so-called *erosion icons*. These are essentially small icons that can be put above each node. The size and color of these erosion icons can then indicate the quantity of code smells for that given node. A controlled experiment has suggested that this gives developers a quicker, more intuitive overview over the distribution of code smells within a project, compared to traditional (i. e., tabular) ways of displaying this information [GKS22]. While LSP does not offer a standardized way of offering code smell information, we can use its diagnostics capability (see section 2.2.2) for the same purpose.

TODO: Show screenshot of erosion icons

2.2 Language Server Protocol

The very basic concepts behind LSP have already been explained in section 1.2.2. The protocol aims to make it easier for IDEs to support more programming languages—specifically, to support language-specific capabilities, which we will go over in section 2.2.2. It was originally developed for Microsoft's editor *Visual Studio Code* (**VSCode**) and was later converted into an open-source specification² (though there are still VSCode-specific extensions to LSP that are not in the official specification today). LSP has found widespread use: An overview page by Microsoft lists at least 269 Language Servers³ (i. e., servers offering support for some programming language) and 61 Language Clients⁴ (i. e., IDEs or development tools).

²Available at <https://github.com/microsoft/language-server-protocol> (*last access: 2024-10-05*).

³<https://microsoft.github.io/language-server-protocol/implementors/servers/> (*last access: 2024-10-05*)

⁴<https://microsoft.github.io/language-server-protocol/implementors/tools/> (*last access: 2024-10-05*)

The current (as of October 6, 2024) version of the protocol is 3.17, with version 3.18 being under active development.

2.2.1 Basics

Messages in the Language Server Protocol are built using JSON-RPC, which uses JSON to encode both requests (consisting of a method name and parameters) and responses (consisting of either a result object or an error object) for procedure calls. Requests include an ID that a response by the server can then reference to match it up with the request it is a reply to. There also so-called *notifications*, which are in essence requests without an ID, intended to send a message server that does not warrant a response [ISO13; Cro06]. In LSP, any message sent between the Language Client and Language Server consist of a header—describing the length and type of the content—followed by the content, which is always a JSON-RPC payload. While the specification does not mandate it, it lists some recommended communication channels on which the protocol messages can be sent, those being stdio (using standard input/output), pipe (using Windows’s pipes), socket (using a socket), and node-ipc (IPC communication over Node.js⁵).

The specification lists all available types for requests and responses as TypeScript interfaces. A sample type definition for the hovering capability, taken from the documentation, can be seen in listing 2.1. From this example, we can also see that locations within documents⁶ are represented as a *Uniform Resource Identifier (URI)* representing the document along with a range (which we have formalized in section 2.1.2). The corresponding method name for this example is `textDocument/hover`.

Listing 2.1: Example specification of request and response objects for the Hover capability.

```
interface HoverParams {
  textDocument: string; /** The text document's URI in string form */
  position: { line: uinteger; character: uinteger; };
}

interface HoverResult {
  value: string;
}
```

The first request from the Language Client to the Language Server always has to be the `initialize` request, including the so-called client capabilities. These specify the capabilities that the Language Client supports. The response from the server will be a response object

⁵<https://nodejs.org> (last access: 2024-10-05)

⁶These documents must always be textual—there is no support for binary files.

that includes the analogous server capabilities. The capability information being sent during this initial handshake is not just a pure list of the names of the corresponding procedure names, but also includes additional details on exactly which parts are supported (or, e. g., which encodings are used), specific to each capability. Afterwards, both parties will then restrict their usage of LSP to the subset of capabilities that both the client and server support.

The end of an LSP session is marked by the Language Client sending the Language Server a shutdown request. After the server confirms the success of the shutdown with a corresponding response, the client should send an exit notification that finally asks the server to quit their process.

For long-running operations, the specification also supports reporting progress on ongoing requests, and cancelling them. The progress reports not only allow indicating the status of the request to the user (e. g., by displaying it in the Language Client's UI), but also allows the server to return partial results to stream responses (e. g., showing the first few references to a variable while the rest are still loading).

The Language Client should also notify the Language Server whenever a document is opened or closed—that way, the server can, for example, start tracking diagnostics in the background as soon as a certain file is opened. There are also notifications related to modifications to the document that the Language Client should send, but these are irrelevant for us because modifying source code is out-of-scope for the LSP integration planned in this thesis.

Apart from describing fundamentals of the protocol like the above, the biggest part of the protocol's specification are the *capabilities*, that is, the JSON-RPC method names and types of the parameters and response objects for each feature that either the Language Client or Language Server can use [LspSpec].

2.2.2 Capabilities

The capabilities I will make use of in SEE can roughly be grouped into the three categories *Navigation*, *Information*, and *Structure*. We will take a detailed look at each relevant capability below, along with how exactly they will be integrated into code cities. This level of detail is justified in the fact that the integration of the capabilities is the main focus of this whole thesis. Note that not all Language Servers support all capabilities—for example, for a language without a hierarchic type system, the capability *type hierarchy* cannot really be sensibly implemented.

Navigation This is the category containing the most capabilities that we can use for this thesis, since there are a lot of ways to navigate from one element within the source code to another. All of these take as input the position in a document, and return any number of locations, where the locations can contain a name, a file, and a range within the file.

All such available capabilities should appear in the context menu of SEE, either upon right-clicking a node, or right-clicking a code element in a code window. Selecting one of these navigation options from the menu should open a menu from which the user can select a single result. If there is only one result to begin with, this step should be skipped. Once a single result has been selected: If the request originates from a code window, the result should be opened and highlighted in that window. If the request instead originates from a code city, the node belonging to that result should be highlighted (e. g., by glowing and having a line pointed to it).

The other important part in SEE where this should be used is when building a city (i. e., when creating the project graph), as these capabilities provides us with the information we need to create (non-hierarchic) edges. Thus, we can create an edge e for each available navigation relation between two nodes, where $\ell(e)$ becomes the type of capability that was used. This is not as trivial as it sounds, since the ranges these capabilities return do not necessarily exactly match the actual ranges of the referenced nodes, so we need to implement some kind of matching algorithm (see section 3.2) .

TODO: More
precise
reference

The following navigation capabilities are available:

- **Call hierarchy:** Returns the incoming/outgoing calls for the symbol at the given location. Since incoming calls are already covered by the references capability, the context menu will only contain an option for showing outgoing calls.
- **Go to declaration:** Returns the declaration location for the symbol at the given location.
- **Go to definition:** Returns the definition location for the symbol at the given location.
 - For this capability, another feature common in IDEs should be implemented in SEE's code windows: Holding `Ctrl`, then clicking on a symbol, directly jumps to the definition of that symbol (or opens the corresponding selection menu, if there is more than one result).
- **Go to implementation:** Returns the definition location for the symbol at the given location.
- **Go to type definition:** Returns the location of the type definition for the symbol at the given location.
- **References:** Returns the location of the references to the symbol at the given location.

- **Type hierarchy:** Returns the sub/supertypes for the symbol at the given location. Since subtypes are already covered by the references capability, the context menu will only contain an option for showing supertypes.

Information Using these capabilities, the user can get information about either the project as a whole, or certain components of it. I have grouped the following capabilities into this category:

- **Diagnostics:** Returns diagnostics for a given file (e. g., warnings or errors). These will be integrated in exactly the same way as the Axivion Suite’s code smells in my bachelor’s thesis were [Gal21], that is:
 - In code cities, we will display erosion icons above affected nodes (see section 2.1.3).
 - In code windows, the corresponding parts of the source code should be highlighted, while hovering over the highlighted parts should reveal the diagnostic’s message and details.

Note that, instead of this being a proper request followed by a server response, this is only the case for the *pull diagnostics* capability, which has been added rather recently . The far more commonly used version is the *push diagnostics* capability, where the Language Server sends out diagnostics for the currently opened files at its own discretion, as notifications (see section 2.2.1)—this makes it difficult to collect this information during city construction, a topic we will explore in section 3.2.

- **Hover:** Returns hover information for a given location. The specification does not specify what exactly this “hover text” should be [LspSpec], but most implementations of Language Servers display the documentation of the hovered element, or the signature if it’s a method, or other helpful associated details. We can simply implement this part of the specification as intended: If the user hovers above an element in a code window, or a node in a code city, we should reveal the hover information in some kind of box near the mouse cursor, hiding it again once the cursor is moved away.
- **Semantic tokens:** Returns semantic tokens for the given file, which are intended for syntax highlighting. Similar to normal (e. g., lexer-based) syntax highlighting, requesting semantic tokens for a document yields a list of tokens containing their positions and a type, where the type can be one that is specified in the protocol (e. g., enum) or one that was previously announced as supported in the client capabilities. IDEs can then render each token type in a different color. An interesting addition to usual syntax highlighting is that each token can also be affected by any number

TODO:
Inconsistent use
of should/will in
this section
TODO: When?

of *token modifiers*, where each modifier may add an additional rendering effect on top of the type-based color. For example, tokens with the static modifier might be rendered in *italics*, while ones with the deprecated modifier could be rendered in ~~strikethrough~~.

SEE currently uses ANTLR⁷-based syntax highlighting, where we need to manually group each parser's token into some categories to determine colors. The added value of the semantic tokens capability here would be ease of use (i. e., no need to manually configure each Language Server) on the one hand, and support for token modifiers (i. e., “extended” syntax highlighting) on the other hand.

Structure This category actually only comprises a single capability—one that we can use to build the hierarchy of the code city's project graph (outlined in section 2.1.2), because it gives us information about the structure of the project. The capability I am talking about here is the *document symbols* capability. Given a document, it will return all symbols present within that file, along with some additional information for each symbol, such as its type or range.

There are actually two different possible kinds of symbols that this capability may return:

1. an array of `SymbolInformation`, which is “a flat list of all symbols” ([LspSpec]) that should not be used to infer a hierarchy. Because of this limitation, if a Language Server is only able to return symbols of this data type, we cannot use it to build code cities. This is an older data type, and instead modern Language Servers should rather return
2. an array of `DocumentSymbol`. This contains a field `children`, which stores `DocumentSymbols` that are contained in this one. Using this property, we can establish a hierarchy and build a code city by recursively enumerating all `DocumentSymbols` and their children for each file, then querying for all relevant information by using the other capabilities outlined above.

Unused Capabilities There are a number of capabilities that I will not implement into SEE. These can be grouped into roughly three categories, based on the reasoning behind them being unused: The first concerns those capabilities that relate to editing only and provide no features related to simply viewing code—as explained in section 1.3, editing code is not part of the goals here and requires additional large-scale preparatory changes to SEE. The second concerns the complex capabilities, that is, ones whose implementation would take a lot of time and effort and thus go beyond the scope of this master's thesis. The third

⁷<https://www.antlr.org> (last access: 2024-06-10)

concerns the niche capabilities that provide only a very marginal benefit, or do so only in rare situations. For these, I also have not deemed the effort worth it to implement them, at least not as part of this thesis.

I will quickly list the contents of all these groups here.

1. Editing capabilities

- **Code Actions:** Allows the programmer to apply refactoring actions to the code, such as importing a referenced library.
- **Completion:** Computes autocomplete items while the user is typing, and applies them when chosen.
- **Formatting:** Applies automatic formatting to a file, or range, of code.
- **Rename:** Executes a project-wide rename of a symbol, which also renames all references to that symbol.
- **Linked Editing Range:** Returns a list of ranges that will be edited upon executing a rename of a symbol, with the purpose of highlighting those ranges during the rename.
- **Signature Help:** Returns signature information at the given cursor information. This may seem relevant for our purposes, but it is actually intended to be shown while editing (e. g., highlighting the active parameter as one types), and its information is given by the hover capability anyway in almost all cases.

2. Niche capabilities

- **Document Color:** Lists all color references in the code (e. g., symbolic references like `Colors.red`) along with their color value in the RGB format.
- **Color Presentation:** Allows users to modify color references in the code by using a color picker.
- **Document Link:** Returns the location of links in the document.
- **Code Lens:** Returns commands that can be shown next to source code, such as the number of implementers of an abstract method.
- **Monikers:** This is a description of what symbols a project imports and which ones it exports, and is intended to make relations between multiple projects possible. As LSP usually only deals with a single project at a time, this is more useful in the *Language Server Index Format* (LSIF), whose specification it also originates from [LsifSpec].

New Term: *LSIF* (Language Server Index Format)

A format which language servers can emit to persist LSP-based information about a software project.

New Term: *DAP* (Debug Adapter Protocol)

A protocol that can be viewed as the analogue to LSP for debuggers, with the goal to make it easier to integrate debuggers into development tools.

3. Complex capabilities

- **Folding Range:** Returns ranges that can be collapsed in the code viewer. For example, the contents of a function could be collapsed, leaving only its signature visible. Due to the way code windows are implemented in SEE, this would increase the complexity of the implementation quite a bit.
- **Inline Value:** In debugging contexts, this supplies the contents of a variable with the purpose of displaying them inline in the Language Client, next to the variable itself. It uses the *Debug Adapter Protocol (DAP)* [DapSpec], which has previously already been integrated into SEE [Roh22], but it would take a lot of refactoring work to make the two implementations compatible with each other.
- **Inlay Hint:** Returns textual hints that can be rendered within the source code. An example would be parameter names that are intended to be shown at the call site, such as in the screenshot in fig. 2.1.
- **Notebook-related capabilities:** These are intended for interactive notebook systems such as Jupyter⁸, which SEE does not currently support.

2.3 Interim Conclusion

In this chapter, we have taken a detailed look at the concepts behind the two topics central to this thesis—namely, the Language Server Protocol and code cities. We have also motivated and laid out the specific ways in which the existing LSP capabilities could be integrated into SEE, including a formalization of the project graph that will become central to section 3.2.

We are now ready to tackle the actual implementation in the next chapter. As a quick overview and recap before then, there is a table of planned capabilities along with their intended use in SEE in table 2.1.

⁸See <https://jupyter.org/> (last access: 2024-10-03)

figures/inlay_hints_example.png

Figure 2.1: An example of inlay hints in the JetBrains IntelliJ IDE. (From <https://www.jetbrains.com/help/idea/inlay-hints.html> (last access: 2024-10-04))

Table 2.1: LSP capabilities that will be integrated into SEE as part of this thesis.

Capability	Code Windows	Code Cities
<i>Call hierarchy</i>	Show incoming/outgoing calls and allow jumping to caller	Generate corresponding edges
<i>Diagnostics</i>	Highlight corresponding code ranges and display details on hover	Display code smell icons [see Gal21]
<i>References</i>	Show references and allow jumping to usage	Generate corresponding edges
<i>Document symbols</i>	—	Generate corresponding nodes and hierarchy
<i>Go to location*</i>	Show locations and allow jumping to them	Generate corresponding edges
<i>Hover</i>	Show hover information when hovering above item	Show hover information when hovering above node
<i>semantic tokens</i>	Extended (“semantic”) syntax highlighting	—
<i>Type hierarchy</i>	Show sub-/supertypes and allow jumping to them	Generate corresponding edges

*This includes the *Go to declaration / definition / implementation / type definition* capabilities.

3

Implementation



TODO!

3.1 Preliminary Changes

TODO!

3.1.1 Specification Cleanup

TODO!

3.1.2 Preparing SEE

TODO!

3.2 Generating Code Cities using LSP

TODO!

3.3 Integrating LSP Functionality into Code Cities

TODO!

3.4 Integrating LSP Functionality into Code Windows

TODO!

3.5 Technical Evaluation

TODO!

3.6 Interim Conclusion

TODO!

4

Evaluation



TODO!

4.1 Plan

TODO!

4.2 Structure

TODO!

4.2.1 Questionnaire

TODO!

4.2.2 Tasks

TODO!

4.3 Results

TODO!

4.4 Threats to Validity

TODO!

4.5 Interim Conclusion

TODO!

5

Conclusion



TODO!

5.1 Limitations

TODO!

5.2 Future Work

TODO!

5.3 The End

TODO!

TODO: Find
better title here

A

List of TODOs



PEN tasks/notes/mistakes for this master's thesis are collected within this appendix. If you see any mistake or empty section not covered by such a note, please tell me. Note that this appendix will only appear in draft versions.

Make sure this is centered!	i
Unfinished section: ()	iii
Unfinished section: ()	v
Convert diagrams to TikZ	1
Use higher quality image.	4
Motivate integration	7
Maybe remove the next part if already mentioned before.	7
In general, reword some stuff from “planning to do” \Rightarrow “done”	7
This is rather vague. Is that alright or do I need to operationalize here?	9
Back this up with sources	11
Refer back here later, noting how LSP can generate code cities for each of these use cases	12
is this correct? And put source here	12
Needs a 3D SEE screenshot, esp. for edges	12
add 2D screenshot	12
Source for 3D/game/VR being easier to navigate/understand	12
link the video	13

A lot of “represents” here	13
Give examples of software projects here to illustrate graph usage?	14
Show screenshot of context menu (pre-LSP)	14
Provide screenshot of editor	14
Show screenshot of erosion icons	15
More precise reference	18
Inconsistent use of should/will in this section	19
When?	19
Unfinished section: Implementation (3)	25
Unfinished section: Preliminary Changes (3.1)	25
Unfinished section: Specification Cleanup (3.1.1)	25
Unfinished section: Preparing SEE (3.1.2)	25
Unfinished section: Generating Code Cities using LSP (3.2)	25
Unfinished section: Integrating LSP Functionality into Code Cities (3.3)	25
Unfinished section: Integrating LSP Functionality into Code Windows (3.4)	26
Unfinished section: Technical Evaluation (3.5)	26
Unfinished section: Interim Conclusion (3.6)	26
Unfinished section: Evaluation (4)	27
Unfinished section: Plan (4.1)	27
Unfinished section: Structure (4.2)	27
Unfinished section: Questionnaire (4.2.1)	27
Unfinished section: Tasks (4.2.2)	27
Unfinished section: Results (4.3)	27
Unfinished section: Threats to Validity (4.4)	28
Unfinished section: Interim Conclusion (4.5)	28
Unfinished section: Conclusion (5)	29
Unfinished section: Limitations (5.1)	29

Unfinished section: Future Work (5.2) 29

Unfinished section: The End (5.3) 29

Find better title here 29

Use higher quality image. 41

B

Glossary

- Graph provider** A component in SEE that produces or transforms a project graph by some means. Can optionally take configuration parameters by users. 14
- Capability** A specific set of features a given Language Server (and Language Client) support. 6, 7, 11, 15–21
- Code City** In the code city metaphor, software components are visualized as buildings within a city, where various metrics of the software are represented visually (e. g., the height of a building could represent the lines of code of the component). 2, 3, 7, 9–12, 14, 16–19, 39
- Code Smell** Certain structures in source code that suggest that a refactoring is in order, such as duplicated code, or a very long method [FB19, pp. 85–87]. 11, 15, 18
- Code Window** A source code viewer in SEE which supports very basic IDE-like functionality. 7, 9, 14, 15, 17, 18, 20
- Polytree** A directed acyclic graph which also has no undirected cycles. 13, 14
- Source Range** The contiguous portion within a source code file that a certain element occupies. In this thesis, we will express it as an ordered pair of a start and end position (s, e) , or, alternatively, as a 4-tuple (s_l, s_c, e_l, e_c) . In the latter representation, the first two elements describe the zero-indexed start line and start character offset, respectively, and the last two describe the corresponding (exclusive) end line and end character offset. 14, 16, 17, 19, 20, 22
- Reflexion Analysis** The process of comparing the architecture and implementation of a software project and finding incongruencies between the two. 11
- Semantic token** An LSP Capability that returns tokens for a document, with the intention that those tokens can be used to apply syntax highlighting to the file. Apart from token types, it also offers token *modifiers*, which can be used to apply additional formatting on top of the colors from the types. 18
- Language Client** A development tool, such as an IDE, that supports LSP and can hence integrate language-specific features into itself using compatible Language Servers. 6, 7, 9, 11, 15, 16, 21

Language Server A locally-running JSON-RPC-based application following the Language Server Protocol that provides language-specific features and aids to the Language Client. 4, 6–9, 11, 15, 16, 18, 19, 39

Unity Editor The main UI of the Unity game engine, in which scenes can be set up, components can be configured, the game itself can be run, etc. Note that it is only used for development purposes, and hence not included within generated builds of a game. 14



Acronyms

DAP A protocol that can be viewed as the analogue to LSP for debuggers, with the goal to make it easier to integrate debuggers into development tools. 21

GXL A file format for graphs, used in SEE for representing dependency and hierarchy graphs of software projects. 3, 13, 14

IDE Editor for source code with features that are useful for development (e. g., highlighting errors). Examples are *Eclipse* or *JETBRAINS IntelliJ*. 3, 5–7, 9–11, 15, 17, 21, 39

JSON-RPC A remote procedure call protocol that uses JSON as its encoding, supporting (among other features) asynchronous calls and notifications. It is used as the base for LSP (even though LSP is technically not a remote protocol). 7, 15, 16

LOC The number of lines in a source code file. 2

LSIF A format which language servers can emit to persist LSP-based information about a software project. 20, *see also* LSP

LSP A protocol which specifies how language servers can provide language-specific features to IDEs, such as hover information, go to definition, or diagnostics. 4–15, 20, 39

SEE An interactive software visualization tool using the code city metaphor in 3D, developed in the Unity game engine at the University of Bremen. 2–4, 7–17, 19–21, 37, 39

UI User Interface 12, 14, 15

URI A string that uniquely identifies some resource. 16

D

Attached Files

Notation	Description
SEE.zip	The build of SEE used for the evaluation.

List of Figures

1.1	A Code City visualized in SEE.	3
1.2	An example of what edges can look like in SEE. TODO: Use higher quality image.	4
1.3	An illustration of how LSP can help simplify IDE development.	5
1.4	An example of the <code>texlab</code> Language Server running in <code>NEOVIM</code>	6
1.6	A code window.	7
1.5	A simplified illustration of how SEE currently gains and uses information about software projects, and how the integration of LSP would change that. .	8
2.1	An example of inlay hints in the JetBrains IntelliJ IDE. (From https://www.jetbrains.com/help/idea/inlay-hints.html (<i>last access: 2024-10-04</i>)) .	23

List of Tables

2.1	LSP capabilities that will be integrated into SEE as part of this thesis.	23
-----	---	----

G

List of Listings

- 2.1. Example specification of request and response objects for the Hover capability. 16



Bibliography

- [Ble22] Moritz Blecker. “Developing a Multi-user Source Code Editor for SEE”. BA thesis. University of Bremen, Jan. 31, 2022 (cit. on p. 7).
- [Cro06] Douglas Crockford. *The Application/Json Media Type for JavaScript Object Notation (JSON)*. Request for Comments RFC 4627. Internet Engineering Task Force, July 2006. 10 pp. (cit. on p. 16).
- [DapSpec] *Debug Adapter Protocol*. Microsoft, May 2022 (cit. on p. 22).
- [FB19] Martin Fowler and Kent Beck. *Refactoring: Improving the Design of Existing Code*. Second edition. The Addison-Wesley Signature Series. Boston Columbus New York San Francisco Amsterdam Cape Town Dubai London Munich: Addison-Wesley, 2019. 418 pp. (cit. on pp. 12, 35).
- [Gal21] Falko Galperin. “Visualizing Code Smells in Code Cities”. BA thesis. University of Bremen, Sept. 27, 2021 (cit. on pp. 11, 15, 19, 23).
- [GKS22] Falko Galperin, Rainer Koschke, and Marcel Steinbeck. “Visualizing Code Smells: Tables or Code Cities? A Controlled Experiment”. In: *2022 Working Conference on Software Visualization (VISSOFT)*. 2022 Working Conference on Software Visualization (VISSOFT). Limassol, Cyprus: IEEE, Oct. 2022, pp. 51–62 (cit. on p. 15).
- [JSO13] JSON-RPC Working Group. *JSON-RPC 2.0 Specification*. Jan. 4, 2013 (cit. on p. 16).
- [LsifSpec] *Language Server Index Format*. Microsoft, Sept. 2021 (cit. on p. 21).
- [LspSpec] *Language Server Protocol*. Microsoft, Oct. 2022 (cit. on pp. 4, 6, 17, 19, 20).
- [McC76] T.J. McCabe. “A Complexity Measure”. In: *IEEE Transactions on Software Engineering* 4 (Dec. 1976), pp. 308–320 (cit. on p. 12).
- [Roh22] Ferdinand Rohlfing. “Debug Adapter Protocol in SEE”. BA thesis. University of Bremen, Aug. 2022 (cit. on p. 22).